CS 540: Introduction to Artificial Intelligence

Homework Assignment #4

Assigned: Friday, April 1 Due: Thursday, April 14

Hand-In Instructions

This assignment includes written problems and programming in Java. The written problems must be done individually but the programming problem can be done either individually or in pairs. Hand in your work electronically to the Moodle dropbox called "HW4 Hand-In". Your answers to each written problem should be turned in as separate pdf files called P1.pdf and P2.pdf and put into a folder called <wisc username>-HW4 For example, if your wisc username is jdoe, then put the pdf files in a folder called jdoe-HW4 If you write out your answers by hand, you'll have to scan your answer and convert the file to pdf. Put your name at the top of the first page in each pdf file. Both people on a team must individually answer and submit their answers to the written problems; no collaboration on these problems is allowed!

For the programming problem, only one person in a pair needs to turn in their code. Put all the files you modified from the provided skeleton code plus any additional java class files that you wrote into the same folder called <wisc username>-HW4 If you have a partner, also put a file called README.txt in this folder that says who your partner is, including both their real name and their wisc username.

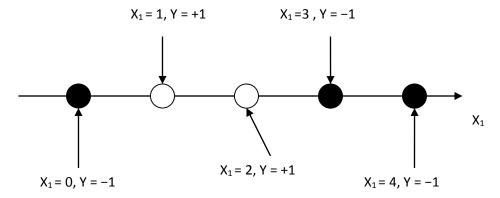
Each person on a team should compress their own folder to create <wisc username>-HW4.zip and copy this one file into the Moodle dropbox. Be sure to verify what you submit!

Late Policy

All assignments are due at 11:59 p.m. on the due date. One (1) day late, defined as a 24-hour period from the deadline (weekday or weekend), will result in 10% of the total points for the assignment deducted. So, for example, if a 100-point assignment is due on a Wednesday and it is handed in between any time on Thursday, 10 points will be deducted. Two (2) days late, 25% off; three (3) days late, 50% off. No homework can be turned in more than three (3) days late. Written questions and program submission have the same deadline. A total of three (3) free late days may be used throughout the semester without penalty. Assignment grading questions must be raised with the instructor or a TA within one week after the assignment is returned.

Problem 1. [10] Support Vector Machines

You are given a dataset with, for each example, a single input feature, X, in \mathfrak{R}^1 , and its desired classification output, $Y \in \{+1, -1\}$, as shown in the figure below. The training set contains three examples in class -1: X = 0, X = 3, and X = 4, and two examples in class +1: X = 1 and X = 2.



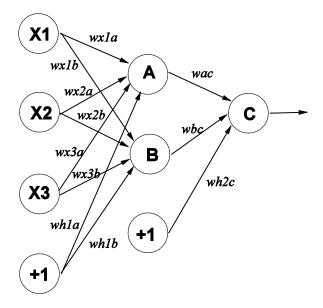
Define a transformation (i.e., a feature mapping function Φ) that maps the data into \Re^2 as follows: $\Phi(X) = (X, X^2)$. Plot the five examples in this 2D space defined by (p, q) where, for each example, p = X, $q = X^2$ and the label at the point is Y.

- (a) [5] Manually construct the maximum-margin separating line by visual inspection of the five points plotted in (a). This will be a line in \Re^2 , which can be parameterized by its normal equation, i.e., $w_1p + w_2q + c = 0$ for appropriate choices of w_1 , w_2 , c. Derive w_1 , w_2 and c.
- (b) [5] Compute the margin for your classifier and circle the support vectors.

Problem 2. [20] Neural Networks

(a) [5] Define a Perceptron with two input units and one output unit that computes the Boolean function $(A \land B) \Leftrightarrow (\neg A \land B)$. Use the definitions in Figure 7.8 in the textbook for the logical connectives \neg , \land and \Leftrightarrow . Let 1 correspond to True and 0 False. Use a step function (LTU) as the activation function. Show the weights and bias values for the network.

(b) [15] Consider a learning task where there are 3 real-valued input units and 1 output unit. You decide to use a 2-layer, feed-forward neural network with 2 hidden units. The activation function is the *sigmoid* function defined in Figure 18.17b in the textbook. Each input unit is connected to every hidden unit, and each hidden unit is connected to the output unit, resulting in the following neural net:



Let the initial weights and biases be given as: wx1a = 0.4, wx1b = -0.2, wx2a = 0.4, wx2b = -0.2, wx3a = 0.4, wx3b = -0.2, wh1a = 0.3, wh1b = 0.3, wac = 0.4, wbc = -0.2, wh2c = 0.3.

- (i) [5] What is the output of nodes A, B and C given a training example with values X1 = 0.3, X2 = 0.8, and X3 = 0.1? Assume that A, B, and C all output real values as computed by their associated *sigmoid* function.
- (ii) [10] Compute *one* (1) step of the Back-propagation algorithm (see Figure 18.24 in the textbook) using the same inputs given in (i) and teacher output C = 1. The error should be computed as the difference between the integer-valued teacher output and the real-valued output of unit C. Using a learning rate of $\alpha = 0.2$. Give your answer as the 11 new weights and biases. Show your work.

Problem 3. [70] Back-propagation for Handwritten Digit Recognition

In this problem you are to write a program that builds a 2-layer, feed-forward neural network and trains it using the back-propagation algorithm. The problem that the neural network will handle is a multi-class classification problem for recognizing handwritten digits. All inputs to the neural network will be numeric. The neural network has one hidden layer only. The network is also fully connected between consecutive layers, meaning each unit, which we'll call a node, in the input layer is connected to all nodes in the hidden layer, and each node in the hidden layer is connected to all the output nodes in the output layer. Each node in the hidden layer and the output layer will also have an extra input from a "bias node" that has constant value +1. So, we can consider both the input layer and the hidden layer as containing one additional node called a bias node. All nodes in the hidden and output layers (except for the bias nodes) should use the ReLU activation function. The initial weights of the network will be randomized. Assuming that input examples (called instances in the code) have m attributes (hence there are m input nodes, not counting the bias node) and we want h nodes (not counting the bias node) in the hidden layer, and o nodes in the output layer, then the total number of weights in the network is (m+1)h between the input and hidden layers, and (h+1)o connecting the hidden and output layers. The number of nodes to be used in the hidden layer will be given as input.

You are required to implement the following two methods from the class NNImpl and one method for the class Node:

```
public class Node{
            public void calculateOutput()
}

public class NNImpl{
            public int calculateOutputForInstance(Instance inst);
            public void train();
}
```

void calculateOutput() calculates the output at a particular node and stores that value in a member variable called outputValue. int calculateOutputForInstance (Instance inst) calculates the output (i.e., the index of the class) for the neural network for a given example (aka instance). void train() trains the neural network using a training set, fixed learning rate, and number of epochs (provided as input to the program).

Dataset

The dataset we will use is called Semeion (https://archive.ics.uci.edu/ml/datasets/Semeion+Handwritten+Digit). It contains 1,593 binary images of size 16 x 16 that each contain one handwritten digit. Your task is to classify each example image as one of the three possible digits: 0, 1 or 2.

Each dataset will begin with a header that describes the dataset: First, there may be several lines starting with "//" that provide a description and comments about the dataset. The line starting with "**" lists the number of output nodes. The line starting with "##" lists the number of attributes, i.e., the number of input values in each instance (in our case, the number of pixels). You can assume that the number of classes will *always* be 3 for this homework because we are

only considering 3-class classification problems. The first output node should output a large value when the instance is determined to be in class 1 (here meaning it is digit 0). The second output node should output a large value when the instance is in class 2 (i.e., digit 1) and, similarly, the third output node corresponds to class 3 (i.e., digit 2). Following these header lines, there will be one line for each instance, containing the values of each attribute followed by the target/teacher values for each output node. For example, if the last 3 values for an instance are: 0 0 1 then this means the instance is the digit 2. We have written the dataset loading part for you according to this format, so do NOT change it.

Implementation Details

We have created four classes to assist your coding, called Instance, Node, NeuralNetworkImpl and NodeWeightPair. Their data members and methods are commented in the skeleton code. We also give an overview of these classes next.

The Instance class has two data members: ArrayList<Double> attributes and ArrayList<Integer> classValues. It is used to represent one instance (aka example) as the name suggests. attributes is a list of all the attributes (in our case binary pixel values) of that instance (all of them are double) and classValues is the class (e.g., 1 0 0 for digit 0) for that instance.

The most important data member of the Node class is int type. It can take the values 0, 1, 2, 3 or 4. Each value represents a particular type of node. The meanings of these values are:

- 0: an input node
- 1: a bias node that is connected to all hidden layer nodes
- 2: a hidden layer node
- 3: a bias node that is connected to all output layer nodes
- 4: an output layer node

Node also has a data member double inputValue that is only relevant if the type is 0. Its value can be updated using the method void setInput(double inputValue). It also has a data member ArrayList < NodeWeightPair > parents, which is a list of all the nodes that are connected to this node from the previous layer (along with the weight connecting these two nodes). This data member is relevant only for types 2 and 4. The neural network is fully connected, which means that all nodes in the input layer (including the bias node) are connected to each node in the hidden layer and, similarly, all nodes in the hidden layer (including the bias node) are connected to the node in the output layer. The output of each node in the output layer is stored in double output Value. You can access this value using the method double getOutput(), which is already implemented. You only need to complete the method void calculateOutput(). This method should calculate the output activation value at the node if its type is 2 or 4. The calculated output should be stored in output Value. The formula for calculating this value is determined by the definition of the ReLU activation function, which is defined as $g(x) = \max(0, x)$ where $x = \sum_{i=1}^{n} w_i x_i$ and x_i are the inputs to the given node, w_i are the corresponding weights, and n is the number of inputs including the bias. When updating weights you'll also use the derivative of the ReLU, defined

as
$$g'(x) = \begin{cases} 0, & \text{if } x \leq 0 \\ 1, & \text{otherwise} \end{cases}$$

NodeWeightPair has two data members, Node and weight. These should be self

explanatory. NNImpl is the class that maintains the whole neural network. It maintains lists of all the input nodes (ArrayList<Node> inputNodes) and the hidden layer nodes (ArrayList<Node> outputNodes). The last node in both the input layer and the hidden layer is the bias node for that layer. Its constructor creates the whole network and maintains the links. So, you do not have to worry about that. As mentioned before, you have to implement two methods here. The data members ArrayList<Instance> trainingSet, double learningRate and int maxEpoch will be required for this. To train the network, implement the back-propagation algorithm given in textbook (Figure 18.24) or in the lecture slides. You can implement it by updating all the weights in the network after each instance (as is done in the algorithm in the textbook). This is the extreme case of Stochastic Gradient Descent. Finally, remember to change the input values of each input layer node (except the bias node) when using each new training instance to train the network.

Classification

Based on the outputs of the output nodes, int calculateOutputForInstance(Instance inst) classifies the instance as the index of the output node with the *maximum* value. For example if one instance has outputs [0.1, 0.2, 0.5], this instance will be classified as digit 2.

Testing

We will test your program on multiple training and testing sets, and the format of testing commands will be:

java HW4 <numHidden> <learnRate> <maxEpoch> <trainFile> <testFile> where trainFile, and testFile are the names of training and testing datasets, respectively. numHidden specifies the number of nodes in the hidden layer (excluding the bias node at the hidden layer). learnRate and maxEpoch are the learning rate and the number of epochs that the network will be trained, respectively. In order to facilitate debugging, we are providing you with sample training data and testing data in the files train1.txt and test1.txt. A sample test command is

```
java HW4 5 0.01 100 train1.txt test1.txt
```

You are NOT responsible for any input or console output. We have written the class HW4 for you, which will load the data and pass it to the method you are implementing. Do NOT modify any IO code. As part of our testing process, we will unzip the files you submit to Moodle, remove any classes, call <code>javac</code> HW4.java to compile your code, and then call the main method, HW4, with parameters of our choosing.

Deliverables

- 1. Hand in your modified version of the code skeleton that includes your implementation of the back-propagation algorithm. *Do not submit any unmodified .java files*. Do include any additional java class files needed to run your program.
- 2. Optionally, submit a file called P3.pdf containing any comments about your program that you would like us to know.